

# CIFAR-10 Image Classification

Cameron Wilson, Brandon Santore

CSCI-335

Rochester Institute of Technology

12/3/2025

# Introduction

airplane	
automobile	
bird	
cat	
deer	
dog	
frog	
horse	
ship	
truck	

# Problem

Given a dataset of CIFAR-10 images, our goal was to design a CNN to be as accurate as possible in classifying the images and compare it to traditional models.

# Background

- We trained three different image-classification models
  - Convolutional Neural Network (CNN)
  - Random Forest
  - Linear Support Vector Machine (SVM)
- CNNs are generally considered to be the best for image classification, so this was our primary model.
- The traditional models were made for comparison

# Hypothesis

Our hypothesis is that out of all of our models our CNN will perform the best as CNNs are commonly thought of as the best models for image classification.

# Methods

# CIFAR-10 Dataset

- We obtained the CIFAR-10 Dataset from the Tensorflow Keras datasets package
- The dataset contains 60,000 images
  - 50,000 training images
  - 10,000 testing images
- The images were 32x32 pixels and colored
- The dataset was divided into 10 classes
  - Airplane, Automobile, Bird, Cat, Deer, Dog, Frog, Horse, Ship, Truck
- The dataset is loaded into Numpy arrays, converted to floats, normalized, and the target arrays are flattened.

# Traditional Models

- We selected Random Forest and Support Vector Classifier as they are fairly good models for multi-classification
- Linear kernel was used for SVC, otherwise the large size of dataset would take too long to train

# Preprocessing

- For the CNN:
  - Pixel values are normalized from [0, 255] to [0, 1]
  - Training data is split: 45,000 images for training and 5,000 for validation
  - One-hot encode the image labels into vectors to allow model to train off of them
- For the SVC and Random Forest:
  - Each image in the dataset is converted to greyscale
  - Features are extracted using Histogram of Oriented Gradients (HOG)

# Histogram of Oriented Gradients

- Analyzes orientation of color gradients in an image
- Extracts rough shapes as features for training



Input image




Histogram of Oriented Gradients

scikit-image team. (n.d.). *Histogram of oriented gradients#*.  
Histogram of Oriented Gradients.  
[https://scikit-image.org/docs/0.25.x/auto\\_examples/features\\_detection/plot\\_hog.html#id2](https://scikit-image.org/docs/0.25.x/auto_examples/features_detection/plot_hog.html#id2)

# CNN Model

- Based on VGG architecture
- Convolution is done in blocks, with multiple passes
- Our model is smaller than most VGG models due to small image size



# Hyperparameters

- Traditional models were not tuned, as we want to use them as a baseline for comparing to our CNN
- Any parameters not specified here were left as default values
- Random Forest:
  - n\_estimators: 100
- SVC:
  - Linear Kernel (used LinearSVC for faster training)
  - C: 1.00
  - max\_iter: 5000
- CNN params were optimized using Adam optimizer
  - Adaptive Moment Estimation
  - Calculate individual learning rates for different parameters
  - Learning rates adapt based on gradients

# Results

# Random Forest Results

- Highest Precision: Airplane
- Lowest Precision: Cat
- Highest Recall: Automobile
- Lowest Recall: Cat
- Highest F1-Score: Airplane and Automobile
- Lowest F1-Score: Cat
- Overall Accuracy: 51%

Precision = TP / (TP + FP)


Recall = TP / (TP + FN)

F1-Score =  $2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$

Random Forest Classification Report:				
	precision	recall	f1-score	support
airplane	0.61	0.61	0.61	1000
automobile	0.57	0.67	0.61	1000
bird	0.46	0.34	0.39	1000
cat	0.35	0.25	0.29	1000
deer	0.41	0.43	0.42	1000
dog	0.40	0.44	0.42	1000
frog	0.51	0.63	0.56	1000
horse	0.57	0.51	0.54	1000
ship	0.57	0.59	0.58	1000
truck	0.59	0.60	0.60	1000
accuracy			0.51	10000
macro avg	0.50	0.51	0.50	10000
weighted avg	0.50	0.51	0.50	10000

# Random Forest Results

- Performed fairly poorly
- Really struggled with animals
- Did fairly well with non-organic objects



# SVC Results

- Highest Precision: Automobile
- Lowest Precision: Cat
- Highest Recall: Frog
- Lowest Recall: Cat
- Highest F1-Score: Truck
- Lowest F1-Score: Cat
- Overall Accuracy: 50%

Precision = TP / (TP + FP)


Recall = TP / (TP + FN)

F1-Score = 2 x (Precision x Recall)/(Precision + Recall)

SVC Classification Report:				
	precision	recall	f1-score	support
airplane	0.55	0.61	0.58	1000
automobile	0.56	0.64	0.60	1000
bird	0.44	0.35	0.39	1000
cat	0.41	0.20	0.27	1000
deer	0.44	0.34	0.38	1000
dog	0.44	0.43	0.43	1000
frog	0.48	0.69	0.57	1000
horse	0.54	0.61	0.57	1000
ship	0.52	0.53	0.52	1000
truck	0.57	0.65	0.61	1000
accuracy			0.50	10000
macro avg	0.49	0.50	0.49	10000
weighted avg	0.49	0.50	0.49	10000

# SVC Results

- Performed fairly poorly
- Struggled with animals as well
- Did okay with non-organic objects



# CNN Results

- Highest Precision: Horse
- Lowest Precision: Frog
- Highest Recall: Frog
- Lowest Recall: Cat
- Highest F1-Score: Ship
- Lowest F1-Score: Cat
- Overall Accuracy: 78%

Precision = TP / (TP + FP)


Recall = TP / (TP + FN)

F1-Score = 2 x (Precision x  
Recall)/(Precision + Recall)

CNN test accuracy: 0.7835				
	precision	recall	f1-score	support
airplane	0.82	0.84	0.83	1000
automobile	0.84	0.95	0.89	1000
bird	0.82	0.65	0.73	1000
cat	0.81	0.40	0.53	1000
deer	0.81	0.71	0.76	1000
dog	0.77	0.69	0.73	1000
frog	0.53	0.97	0.69	1000
horse	0.90	0.82	0.86	1000
ship	0.89	0.91	0.90	1000
truck	0.85	0.88	0.87	1000
accuracy			0.78	10000
macro avg	0.81	0.78	0.78	10000
weighted avg	0.81	0.78	0.78	10000

# CNN Results

- Possible overfit with frog classification
- Low cat accuracy
- Fairly good results



# Conclusions

# Interpretations

- The CNN model is a very accurate model for image classification.
- All of our models had a tendency to predict “frog” on the images, with CNN doing it the most. This is most likely due to some similarities in features between the image classes.
  - It is interesting how the model falsely predicted many classes to be “frog” but did not falsely predict “frog” to be those classes.
- All of our models performed poorly with the “cat” classification. “Cat” had the lowest recall and f1-score in every model.

# Comparison

- Our CNN model performed very well, and predictably better than the traditional models
- The traditional models both performed similarly to each other, with little to noticeable differences
  - The SVC performed slightly worse with non-organic objects

# Limitations

- Computing power was a huge limitation on our model complexity
- The dataset itself could be giving us worse results. Many of the image classes have similar color schemes and shapes, which we believe to be a reason why the model is predicting “frog” very often.
  - Many of the animal classes have similar color schemes due to the tendency of animals being found in nature

# Future Work

- Could try testing model on a different dataset of the same image classes.
- More resources would lead to being able to use RBF kernel for SVC and potentially a better comparison
- The complexity of our CNN could also be increased with more computing power
- Increase resolution of images or add more data augmentation.



# Any Questions?